

This mini-project is based on the K-Means exercise from ‘R in Action’

Go here for the original blog post and solutions: <http://www.r-bloggers.com/k-means-clustering-from-r-in-action/>

Exercise 0:

Install these packages if you don’t have them already:

```
install.packages(c("cluster", "rattle.data", "NbClust"))
```

Now load the data and look at the first few rows

```
data(wine, package="rattle.data")
head(wine)
```

##	Type	Alcohol	Malic	Ash	Alcalinity	Magnesium	Phenols	Flavanoids
## 1	1	14.23	1.71	2.43	15.6	127	2.80	3.06
## 2	1	13.20	1.78	2.14	11.2	100	2.65	2.76
## 3	1	13.16	2.36	2.67	18.6	101	2.80	3.24
## 4	1	14.37	1.95	2.50	16.8	113	3.85	3.49
## 5	1	13.24	2.59	2.87	21.0	118	2.80	2.69
## 6	1	14.20	1.76	2.45	15.2	112	3.27	3.39
##	Nonflavanoids	Proanthocyanins	Color	Hue	Dilution	Proline		
## 1	0.28		2.29	5.64	1.04	3.92	1065	
## 2	0.26		1.28	4.38	1.05	3.40	1050	
## 3	0.30		2.81	5.68	1.03	3.17	1185	
## 4	0.24		2.18	7.80	0.86	3.45	1480	
## 5	0.39		1.82	4.32	1.04	2.93	735	
## 6	0.34		1.97	6.75	1.05	2.85	1450	

Exercise 1:


Remove the first column from the data and scale it using the scale() function

```
df = scale(wine[-1])
```

Now we'd like to cluster the data using K-Means. How do we decide how many clusters to use if you don't know that already? We'll try two methods.

Method 1: A plot of the total within-groups sums of squares against the number of clusters in a K-means solution can be helpful. A bend in the graph can suggest the appropriate number of clusters.

```
wssplot <- function(data, nc=15, seed=1234){  
  wss <- (nrow(data)-1)*sum(apply(data,2,var))  
  for (i in 2:nc){  
    set.seed(seed)  
    wss[i] <- sum(kmeans(data, centers=i)$withinss)}  
  
  plot(1:nc, wss, type="b", xlab="Number of Clusters",  
       ylab="Within groups sum of squares")  
}  
  
wssplot(df)
```



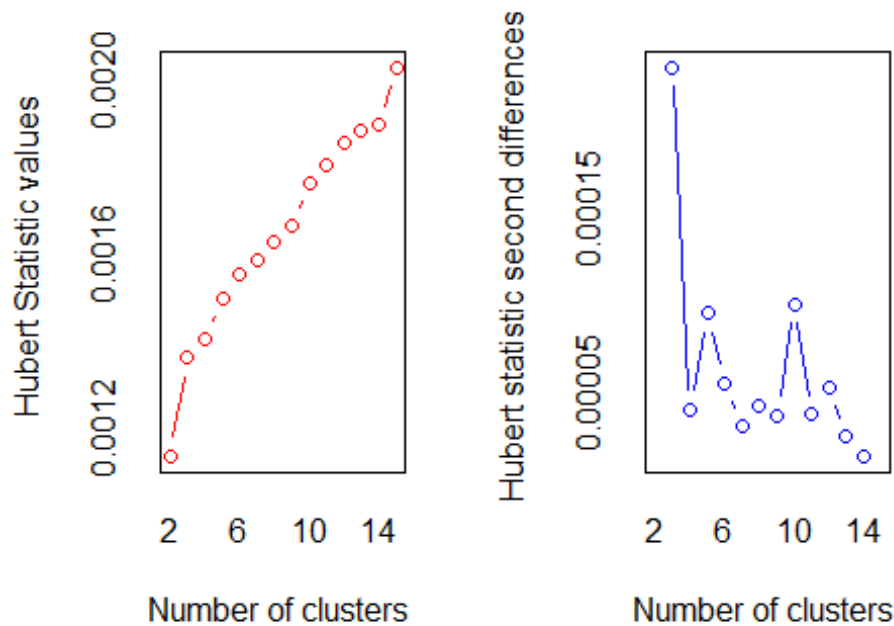
Exercise 2:

- * How many clusters does this method suggest?
- * Why does this method work? What's the intuition behind it?
- * Look at the code for `wssplot()` and figure out how it works

This method suggests 3 clusters. The function for `wssplot` iterates the `kmeans` function across all columns within the `df` and extracts the `$withinss` value. Then, this figure is applied to each column and plotted.


Method 2: Use the `NbClust` library, which runs many experiments and gives a distribution of potential number of clusters.

```
library(NbClust)
set.seed(1234)
nc <- NbClust(df, min.nc=2, max.nc=15, method="kmeans")
```



```
## *** : The Hubert index is a graphical method of determining the number of
clusters.
##           In the plot of Hubert index, we seek a significant knee
that corresponds to a
##           significant increase of the value of the measure i.e the
significant peak in Hubert
```


```
##          index second differences plot.
##
```



```
## *** : The D index is a graphical method of determining the number of
## clusters.
##          In the plot of D index, we seek a significant knee (the
##          significant peak in Dindex
##          second differences plot) that corresponds to a significant
##          increase of the value of
##          the measure.
##
## *****
## * Among all indices:
## * 4 proposed 2 as the best number of clusters
## * 15 proposed 3 as the best number of clusters
## * 1 proposed 10 as the best number of clusters
## * 1 proposed 12 as the best number of clusters
## * 1 proposed 14 as the best number of clusters
## * 1 proposed 15 as the best number of clusters
##
##          ***** Conclusion *****
##
## * According to the majority rule, the best number of clusters is 3
##
## *****
```

```
barplot(table(nc$Best.n[1,]),
        xlab="Numer of Clusters", ylab="Number of Criteria",
        main="Number of Clusters Chosen by 26 Criteria")
```

er of Clusters Chosen by :



Exercise 3:

How many clusters does this method suggest?

This method suggested 3 clusters according to majority rule from 15 indices.

Exercise 4:

Once you've picked the number of clusters, run k-means using this number of clusters. Output the result of calling `kmeans()` into a variable `fit.km`

```
fit.km <- kmeans(df, 3, nstart=25)
```

Now we want to evaluate how well this clustering does.

Exercise 5:

*Using the table() function, show how the clusters in fit.km\$clusters # compares to the actual wine types in wine\$Type.

*Would you consider this a good clustering?

```
wine_val = table(wine$Type, fit.km$cluster)
```

After reviewing the attached link, my numbers are in reverse order, but the actual figures are correct. After reviewing the table, 6 predictions were incorrect, mainly in the second cluster.

Exercise 6:

* Visualize these clusters using function clusplot() from the cluster library

* Would you consider this a good clustering?

```
library(cluster)
```

```
## Warning: package 'cluster' was built under R version 3.4.4
```

```
clusplot(df, fit.km$cluster, main = "Cusplot")
```

